


OPT-13B ShareGPT validation on one A100 40GB

 PagedAttention SOSP 2023  Our vLLM 0.25.1 run

Normalized latency (s/token, log)


(a) Published curve + our point



Normalized latency (s/token)

(b) Same-model result at 1 req/s

Mean gap: -30.6% | Local P99: 0.0349
100/100 completed | Waiting: 0 | Offload: disabled



Matched: model, GPU, FP16, ShareGPT, Poisson 1 req/s, 2K context, metric. Different: vLLM version and trace duration.